



A large, thick, black L-shaped frame surrounds the central text. It consists of a vertical bar on the left and a horizontal bar at the top, meeting at a corner in the upper-left. Another vertical bar is on the right, and a horizontal bar is at the bottom, meeting at a corner in the lower-right.

TRANSPORT LAYER SERVICES AND PROTOCOLS

Dr Chathu Ranaweera
Senior Lecturer in Computer Networks

Transport vs. Network Layer Services and Protocols



© Tracy Smith
www.flicker.com

Household analogy:

Three kids in an Australian house sending letters to three kids in a Belgium house:

- hosts = houses
- processes = kids
- app messages = letters in envelopes

Transport vs. Network Layer Services and Protocols

- **Network layer** : logical communication between *hosts*
- **Transport layer**: logical communication between *processes*
 - relies on, enhances, network layer services


household analogy:

Three kids in an Australian house sending letters to three kids in a Belgium house:

- hosts = houses
- processes = kids
- app messages = letters in envelopes

Two principal Internet transport protocols

- **TCP:** Transmission Control Protocol
 - reliable, in-order delivery
 - congestion control
 - flow control
 - connection setup
- **UDP:** User Datagram Protocol
 - unreliable, unordered delivery
 - no-frills extension of “best-effort” IP
- Services not available (if not supported via network layer):
 - delay guarantees
 - bandwidth guarantees



A thick black L-shaped frame is positioned around the text. It starts at the top-left, goes right, then down, then right again, and finally down to the bottom-right corner.

MULTIPLEXING AND DEMULTIPLEXING


Multiplexing/demultiplexing

multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)


demultiplexing at receiver:

use header info to deliver received segments to correct socket




How demultiplexing works in Host

- Host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries one transport-layer segment
 - each segment has source, destination port number
- Host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format

A thick black L-shaped frame is positioned around the text. It starts at the top-left, goes right, then down, then right again, and finally down to the bottom-right corner.

CONNECTIONLESS AND CONNECTION ORIENTED MULTIPLEXING AND DEMULTIPLEXING

Connectionless demultiplexing

- When creating socket, must specify *host-local* port #
- When creating datagram to send into UDP socket, must specify
 - destination IP address
 - destination port #


when receiving host receives *UDP* segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



IP/UDP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at receiving host


Connectionless demultiplexing: an example



Connection-oriented demultiplexing

- TCP socket identified by **4-tuple**:
 - source IP address
 - source port number
 - destination IP address
 - destination port number
- demux: receiver uses *all four values (4-tuple)* to direct segment to appropriate socket
- server may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
 - each socket associated with a different connecting client


Connection-oriented demultiplexing: example



Three segments, all destined to IP address: B,
destination port: 80 are demultiplexed to *different* sockets

What did we learn?

- Multiplexing, demultiplexing: based on segment, datagram header field values
- **UDP:** demultiplexing using destination port number (only)
- **TCP:** demultiplexing using 4-tuple: source and destination IP addresses, and port numbers
- Multiplexing/demultiplexing happen at *all* layers

A large, thick, black L-shaped frame surrounds the central text. It consists of a vertical bar on the left and a horizontal bar at the top, meeting at a corner in the top-left. Another vertical bar is on the right, and a horizontal bar is at the bottom, meeting at a corner in the bottom-right.

CONNECTIONLESS TRANSPORT LAYER PROTOCOL

Dr Chathu Ranaweera
Senior Lecturer in Computer Networks

UDP: User Datagram Protocol

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
 - lost
 - delivered out-of-order to app
- *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others

Why is there a UDP?

- no connection establishment (which can add RTT delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control
 - UDP can blast away as fast as desired!
 - can function in the face of congestion

UDP: User Datagram Protocol

- UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP (Simple Network Management Protocol)
 - HTTP/3 : QUIC (Quick UDP Internet Connections)
- If reliable transfer needed over UDP (e.g., HTTP/3):
 - Add needed reliability at application layer
 - Add congestion control at application layer

UDP: User Datagram Protocol [RFC 768]

INTERNET STANDARD

RFC 768

J. Postel

ISI

28 August 1980


User Datagram Protocol

Introduction


This User Datagram Protocol (UDP) is defined to make available a datagram mode of packet-switched computer communication in the environment of an interconnected set of computer networks. This protocol assumes that the Internet Protocol (IP) [[1](#)] is used as the underlying protocol.

This protocol provides a procedure for application programs to send messages to other programs with a minimum of protocol mechanism. The protocol is transaction oriented, and delivery and duplicate protection are not guaranteed. Applications requiring ordered reliable delivery of streams of data should use the Transmission Control Protocol (TCP) [[2](#)].

Format




UDP: Transport Layer Actions



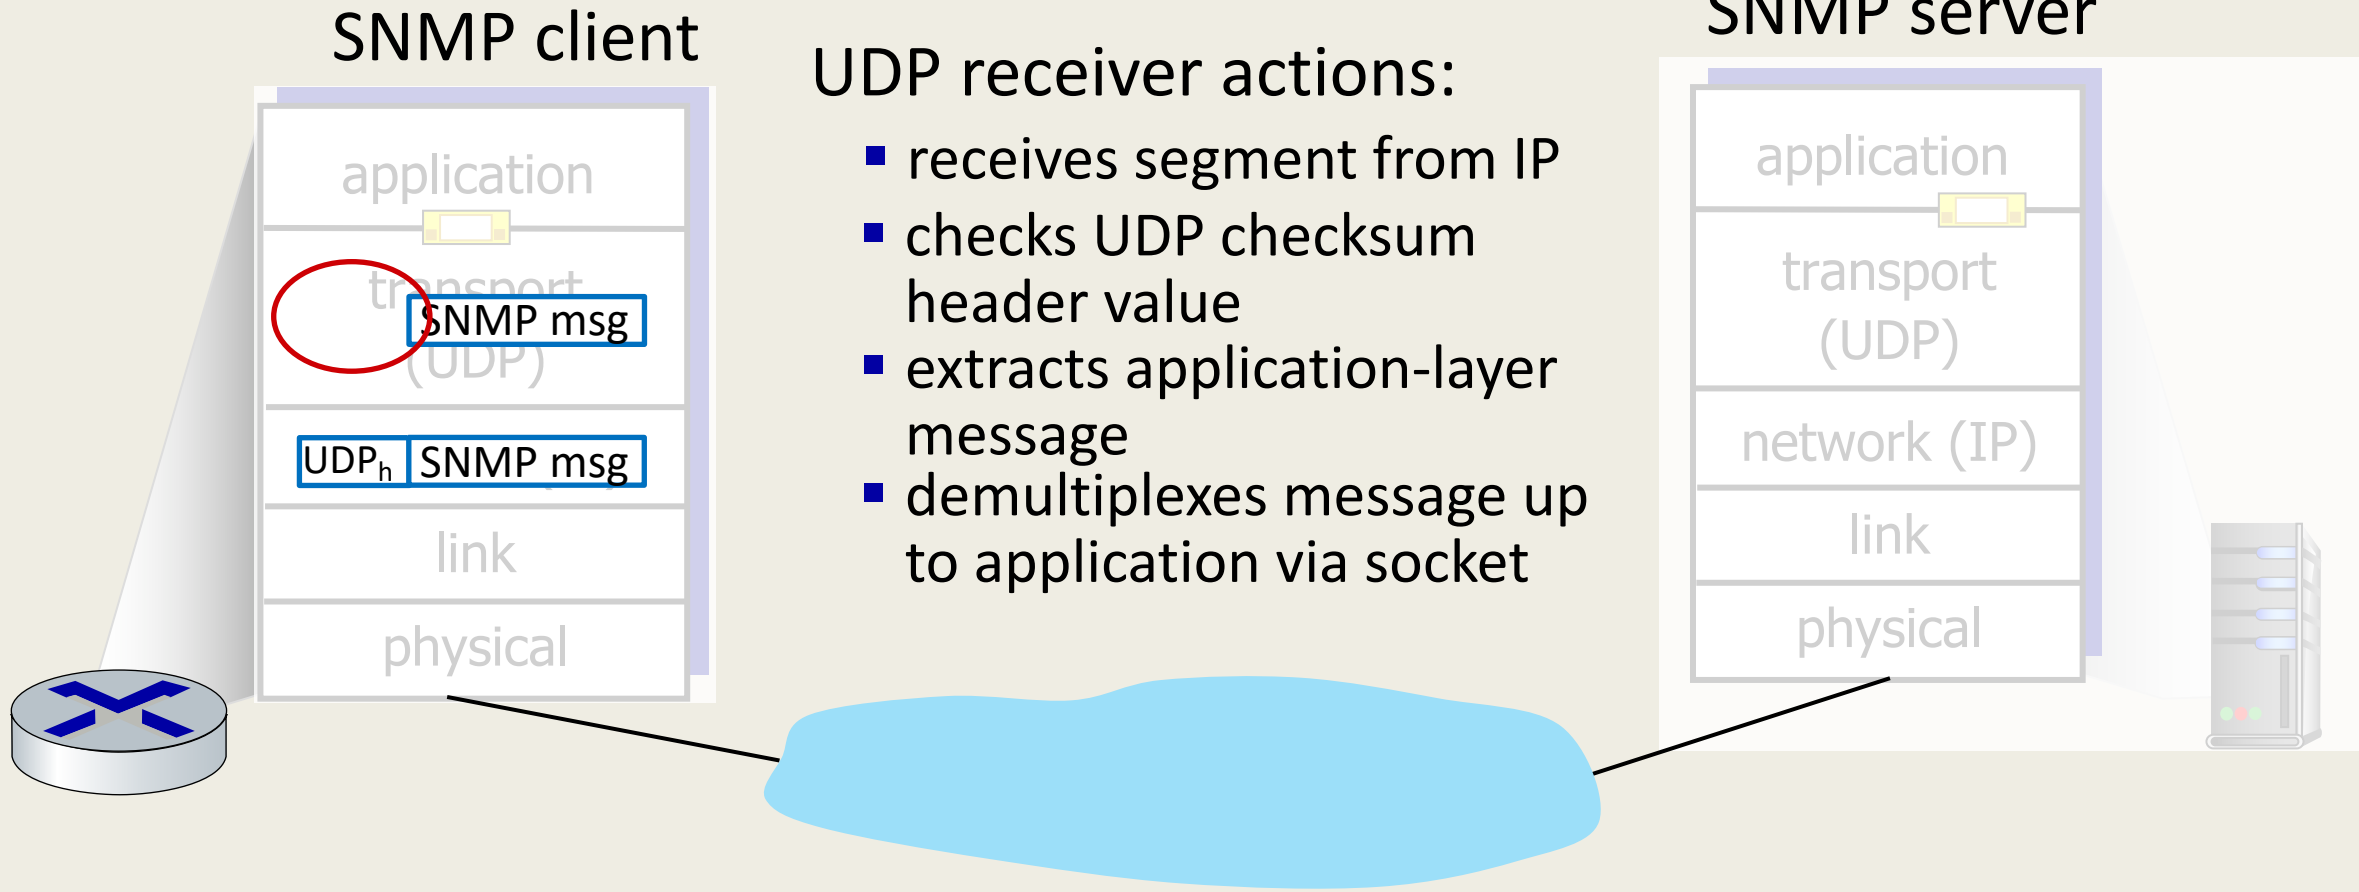
UDP sender actions:

- is passed an application-layer message
- determines UDP segment header fields values
- creates UDP segment
- passes segment to IP

SNMP server




UDP: Transport Layer Actions




UDP : Snapshot

- “no frills” protocol:
 - segments may be lost, delivered out of order
 - best effort service: “send and hope for the best”
- UDP has its plusses:
 - no setup/handshaking needed (no RTT incurred)
 - can function when network service is compromised
 - helps with reliability (checksum)
- build additional functionality on top of UDP in application layer (e.g., HTTP/3)

A diagram showing a UDP segment. It consists of a large black L-shaped bracket on the left side, opening to the right. A second, smaller black L-shaped bracket is on the right side, opening to the left. These two brackets frame the text "UDP SEGMENT" in the center. The text is in a bold, black, sans-serif font.


UDP SEGMENT

UDP segment header



UDP checksum

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment



UDP checksum

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

sender:

- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- **checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - Not equal - error detected
 - Equal - no error detected. *But maybe errors nonetheless?* More later

Internet checksum: an example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Internet checksum: weak protection!

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1


Internet checksum: weak protection!

example: add two 16-bit integers

		1	1	1	0	0	1	0	0	0	1	1	0	0	1	0	0
		1	1	0	1	0	1	1	1	0	1	0	1	0	1	1	1
Wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
Sum after Wraparound		1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
Checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1
		1	1	1	0	0	1	0	0	0	1	1	0	0	1	0	1
		1	1	0	1	0	1	1	1	0	1	0	1	0	1	0	0
Wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
Sum after Wraparound		1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
Checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Bits have changed due to errors

Even though bits have changed, *no* change in checksum!



TCP






TCP SEGMENT STRUCTURE



TCP segment structure





SEQUENCE AND ACKNOWLEDGMENT NUMBERS



TCP sequence numbers, ACKs

Sequence numbers:

- byte stream “number” of first byte in segment’s data

Acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender


source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer




outgoing segment from receiver

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

TCP sequence numbers, ACKs



simple telnet scenario

A large, thick, black L-shaped frame surrounds the central text. It consists of a vertical bar on the left and a horizontal bar at the top, meeting at a corner in the top-left. Another vertical bar is on the right, and another horizontal bar is at the bottom, meeting at a corner in the bottom-right.

ESTABLISHING AND CLOSING TCP CONNECTIONS

Dr Chathu Ranaweera
Senior Lecturer in Computer Networks


Control Flags

- URG: Urgent Pointer - is used to identify incoming data as “urgent”
- ACK: Acknowledgement – is used to indicate the acknowledgement of the receipt of data and the acknowledgement number field should be examined.
- PHS: Push – is used to indicate to push currently -buffered data to the application at the receiving side
- RST: Reset – is used to indicate TCP hasn’t been able to properly recover from misinformed segments.
- SYN: Synchronise – is used when first establishing a TCP connection
- FIN: Finish – is used to indicate that the connection can be closed

Establishing a TCP Connection

Host A :Sender

Host B :Receiver




Three Way Handshake

Closing a TCP Connection

Host A :Sender

Host B :Receiver



Four Way Handshake



TCP RETRANSMISSION



TCP round trip time, timeout

Q: How to set TCP timeout value?

- longer than RTT, but RTT varies!
- *too short*: premature timeout, unnecessary retransmissions
- *too long*: slow reaction to segment loss

Q: How to estimate RTT?

- *SampleRTT*: measured time from segment transmission until ACK receipt
 - ignore retransmissions
- *SampleRTT* will vary, want estimated RTT “smoother”
 - average several *recent* measurements, not just current *SampleRTT*

TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average (EWMA)
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$

TCP round trip time, timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
 - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑
estimated RTT

↑
“safety margin”

- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

TCP Sender (simplified)

event: data received from application

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
 - think of timer as for oldest unACKed segment
 - expiration interval: **TimeoutInterval**


event: timeout

- retransmit segment that caused timeout
- restart timer

event: ACK received

- if ACK acknowledges previously unACKed segments
 - update what is known to be ACKed
 - start timer if there are still unACKed segments

TCP Receiver: ACK generation [RFC 5681]







TCP RETRANSMISSION SCENARIOS



TCP: retransmission scenarios




lost ACK scenario



premature timeout

TCP: retransmission scenarios



TCP fast retransmit

TCP fast retransmit

if sender receives 3 additional ACKs for same data (“triple duplicate ACKs”), resend unACKed segment with smallest seq #

- likely that unACKed segment lost, so don't wait for timeout



Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. So retransmit!

